

## AI and ML task 7

### importing libraries

```
In [5]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report

In [7]: data = pd.read_csv('breast-cancer.csv')

In [9]: data

Out[9]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...	radius_worst	texture_worst	perimeter_worst	area_worst	smoothness_worst
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30010	0.14710	...	25.380	17.33	184.60	2019.0	0.16220
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.08690	0.07017	...	24.990	23.41	158.80	1956.0	0.12380
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.19740	0.12790	...	23.570	25.53	152.50	1709.0	0.14440
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.24140	0.10520	...	14.910	26.50	98.87	567.7	0.20980
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.19800	0.10430	...	22.540	16.67	152.20	1575.0	0.13740
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
564	926424	M	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	0.13890	...	25.450	26.40	166.10	2027.0	0.14100
565	926682	M	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	0.09791	...	23.690	38.25	155.00	1731.0	0.11660
566	926954	M	16.60	28.08	108.30	858.1	0.08455	0.10230	0.09251	0.05302	...	18.980	34.12	126.70	1124.0	0.11390
567	927241	M	20.60	29.33	140.10	1265.0	0.11780	0.27700	0.35140	0.15200	...	25.740	39.42	184.60	1821.0	0.16500
568	92751	B	7.76	24.54	47.92	181.0	0.05263	0.04362	0.00000	0.00000	...	9.456	30.37	59.16	268.6	0.08996

569 rows × 32 columns

```
In [11]: data.shape
Out[11]: (569, 32)

In [13]: data.head()

Out[13]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	...	radius_worst	texture_worst	perimeter_worst	area_worst	smoothness_worst
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30010	0.14710	...	25.38	17.33	184.60	2019.0	0.1622
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017	...	24.99	23.41	158.80	1956.0	0.1238
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790	...	23.57	25.53	152.50	1709.0	0.1444
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520	...	14.91	26.50	98.87	567.7	0.2098
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430	...	22.54	16.67	152.20	1575.0	0.1374

5 rows × 32 columns

```
In [15]: data.describe()

Out[15]:
```

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry_mean	...	radius_worst	texture_worst	perimeter_worst	area_worst	smoothness_worst
count	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	...	569.000000	569.000000	569.000000	569.000000	569.000000
mean	3.037183e+07	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799	0.048919	0.181162	...	16.269190	25.677223	107.261213	880.583128	0.137400
std	1.250206e+08	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720	0.038803	0.027414	...	4.833242	6.146258	33.602542	569.356993	0.014400
min	8.670000e+03	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000	0.000000	0.106000	...	7.930000	12.020000	50.410000	185.200000	0.089960
50%	6.692180e+05	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560	0.020310	0.161900	...	13.010000	21.080000	84.110000	515.300000	0.113900
25%	9.060240e+05	13.370000	18.840000	86.240000	551.100000	0.095870	0.026390	0.081540	0.033500	0.179200	...	14.970000	25.410000	97.660000	686.500000	0.165000
75%	8.813129e+06	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130700	0.074000	0.195700	...	18.790000	29.720000	125.400000	1084.000000	0.209800
max	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800	0.201200	0.304000	...	36.040000	49.540000	251.200000	4254.000000	0.277000

8 rows × 31 columns

```
In [17]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
 #   Column                Non-Null Count  Dtype
---  --
 0   id                    569 non-null    int64
 1   diagnosis             569 non-null    object
 2   radius_mean           569 non-null    float64
 3   texture_mean          569 non-null    float64
 4   perimeter_mean        569 non-null    float64
 5   area_mean             569 non-null    float64
 6   smoothness_mean       569 non-null    float64
 7   compactness_mean      569 non-null    float64
 8   concavity_mean        569 non-null    float64
 9   concave points_mean   569 non-null    float64
10   symmetry_mean         569 non-null    float64
11   fractal_dimension_mean 569 non-null    float64
12   radius_se             569 non-null    float64
13   texture_se            569 non-null    float64
14   perimeter_se          569 non-null    float64
15   area_se              569 non-null    float64
16   smoothness_se        569 non-null    float64
17   compactness_se       569 non-null    float64
18   concavity_se         569 non-null    float64
19   concave points_se     569 non-null    float64
20   symmetry_se          569 non-null    float64
21   fractal_dimension_se  569 non-null    float64
22   radius_worst          569 non-null    float64
23   texture_worst         569 non-null    float64
24   perimeter_worst       569 non-null    float64
25   area_worst           569 non-null    float64
26   smoothness_worst     569 non-null    float64
27   compactness_worst    569 non-null    float64
28   concavity_worst      569 non-null    float64
29   concave points_worst  569 non-null    float64
30   symmetry_worst       569 non-null    float64
31   fractal_dimension_worst 569 non-null    float64
dtypes: float64(30), int64(1), object(1)
memory usage: 142.4+ KB

In [19]: data.isnull().sum()

Out[19]:
id                0
diagnosis         0
radius_mean      0
texture_mean     0
perimeter_mean  0
area_mean        0
smoothness_mean 0
compactness_mean 0
concavity_mean   0
concave points_mean 0
symmetry_mean    0
fractal_dimension_mean 0
radius_se        0
texture_se       0
perimeter_se     0
area_se          0
smoothness_se    0
compactness_se   0
concavity_se     0
concave points_se 0
symmetry_se      0
fractal_dimension_se 0
radius_worst     0
texture_worst    0
perimeter_worst  0
area_worst       0
smoothness_worst 0
compactness_worst 0
concavity_worst  0
concave points_worst 0
symmetry_worst   0
fractal_dimension_worst 0
dtype: int64
```

### preparing dataset for binary classification

```
In [22]: le = LabelEncoder()
data['diagnosis'] = le.fit_transform(data['diagnosis'])

In [24]: X = data.drop('diagnosis', axis=1)
y = data['diagnosis']

In [26]: scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

In [28]: X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42
)
```

### train svm with linear and rbf kernel

```
In [31]: # svm with linear kernel
svm_linear = SVC(kernel='linear', C=1.0, random_state=42)
svm_linear.fit(X_train, y_train)
y_pred_linear = svm_linear.predict(X_test)

In [33]: # Train SVM with RBF kernel
svm_rbf = SVC(kernel='rbf', C=1.0, gamma='scale', random_state=42)
svm_rbf.fit(X_train, y_train)
y_pred_rbf = svm_rbf.predict(X_test)

In [35]: # Evaluation
print("SVM with Linear Kernel:")
print("Accuracy:", accuracy_score(y_test, y_pred_linear))
print(classification_report(y_test, y_pred_linear))

print("\nSVM with RBF Kernel:")
print("Accuracy:", accuracy_score(y_test, y_pred_rbf))
print(classification_report(y_test, y_pred_rbf))

SVM with Linear Kernel:
Accuracy: 0.95610350871193
precision    recall  f1-score   support

   0      0.97      0.96      0.96      71
   1      0.93      0.95      0.94      43

 accuracy      0.95      0.96      0.95     114
 macro avg      0.95      0.96      0.95     114
weighted avg      0.96      0.96      0.96     114

SVM with RBF Kernel:
Accuracy: 0.9824561403508771
precision    recall  f1-score   support

   0      0.97      1.00      0.99      71
   1      1.00      0.95      0.98      43

 accuracy      0.99      0.98      0.98     114
 macro avg      0.99      0.98      0.98     114
weighted avg      0.98      0.98      0.98     114
```

### Visualize decision boundary using 2D data.

```
In [40]: from sklearn.decomposition import PCA

In [42]: # Reduce features to 2D using PCA
pca = PCA(n_components=2)
X_train_2D = pca.fit_transform(X_train)
X_test_2D = pca.transform(X_test)



In [44]: # Train new SVMs on 2D data
svm_linear_2D = SVC(kernel='linear', C=1.0).fit(X_train_2D, y_train)
svm_rbf_2D = SVC(kernel='rbf', C=1.0, gamma='scale').fit(X_train_2D, y_train)

def plot_decision_boundary(model, X, y, title):
    h = 0.02 # step size in the mesh
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                        np.arange(y_min, y_max, h))

    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    plt.figure(figsize=(6, 5))
    plt.contourf(xx, yy, Z, alpha=0.3, cmap=plt.cm.coolwarm)
    plt.scatter(X[:, 0], X[:, 1], c=y, edgecolors='k', cmap=plt.cm.coolwarm)
    plt.title(title)
    plt.xlabel("Principal Component 1")
    plt.ylabel("Principal Component 2")
    plt.tight_layout()
    plt.show()

In [46]: plot_decision_boundary(svm_linear_2D, X_train_2D, y_train, "SVM Linear Kernel (2D PCA)")
plot_decision_boundary(svm_rbf_2D, X_train_2D, y_train, "SVM RBF Kernel (2D PCA)")
```



### Tune hyperparameters like C and gamma.

```
In [54]: from sklearn.model_selection import GridSearchCV

In [56]: # Define parameter grids
param_grid_linear = {
    'C': [0.01, 0.1, 1, 10, 100]
}

param_grid_rbf = {
    'C': [0.1, 1, 10, 100],
    'gamma': [0.01, 0.1, 1, 'scale', 'auto']
}

In [58]: # SVM with linear kernel
grid_linear = GridSearchCV(SVC(kernel='linear'), param_grid_linear, cv=5, scoring='accuracy')
grid_linear.fit(X_train, y_train)

Out[58]:
> GridSearchCV
> best_estimator_: SVC

In [60]: # SVM with RBF kernel
grid_rbf = GridSearchCV(SVC(kernel='rbf'), param_grid_rbf, cv=5, scoring='accuracy')
grid_rbf.fit(X_train, y_train)

Out[60]:
> GridSearchCV
> best_estimator_: SVC

In [62]: print("Best Linear SVM Parameters:", grid_linear.best_params_)
print("Best Linear SVM Score:", grid_linear.best_score_)

print("\nBest RBF SVM Parameters:", grid_rbf.best_params_)
print("Best RBF SVM Score:", grid_rbf.best_score_)

Best Linear SVM Parameters: {'C': 0.1}
Best Linear SVM Score: 0.9736263736263737

Best RBF SVM Parameters: {'C': 1, 'gamma': 'scale'}
Best RBF SVM Score: 0.9736263736263737

In [66]: # Evaluate on test set
best_linear = grid_linear.best_estimator_
best_rbf = grid_rbf.best_estimator_

In [68]: from sklearn.metrics import classification_report

print("\nTest Set Performance - Best Linear SVM:")
print(classification_report(y_test, best_linear.predict(X_test)))

print("\nTest Set Performance - Best RBF SVM:")
print(classification_report(y_test, best_rbf.predict(X_test)))

Test Set Performance - Best Linear SVM:
precision    recall  f1-score   support

   0      0.97      1.00      0.99      71
   1      1.00      0.95      0.98      43

 accuracy      0.99      0.98      0.98     114
 macro avg      0.99      0.98      0.98     114
weighted avg      0.98      0.98      0.98     114

Test Set Performance - Best RBF SVM:
precision    recall  f1-score   support

   0      0.97      1.00      0.99      71
   1      1.00      0.95      0.98      43

 accuracy      0.99      0.98      0.98     114
 macro avg      0.99      0.98      0.98     114
weighted avg      0.98      0.98      0.98     114
```

