

```
In [3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, classification_report

In [4]: data = pd.read_csv('Iris.csv')

In [5]: data

Out[5]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
	0	1	5.1	3.5	1.4	0.2 Iris-setosa
	1	2	4.9	3.0	1.4	0.2 Iris-setosa
	2	3	4.7	3.2	1.3	0.2 Iris-setosa
	3	4	4.6	3.1	1.5	0.2 Iris-setosa
	4	5	5.0	3.6	1.4	0.2 Iris-setosa
...
	145	146	6.7	3.0	5.2	2.3 Iris-virginica
	146	147	6.3	2.5	5.0	1.9 Iris-virginica
	147	148	6.5	3.0	5.2	2.0 Iris-virginica
	148	149	6.2	3.4	5.4	2.3 Iris-virginica
	149	150	5.9	3.0	5.1	1.8 Iris-virginica

150 rows x 6 columns

```
In [9]: data.shape

Out[9]: (150, 6)

In [11]: data.head()

Out[11]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
	0	1	5.1	3.5	1.4	0.2 Iris-setosa
	1	2	4.9	3.0	1.4	0.2 Iris-setosa
	2	3	4.7	3.2	1.3	0.2 Iris-setosa
	3	4	4.6	3.1	1.5	0.2 Iris-setosa
	4	5	5.0	3.6	1.4	0.2 Iris-setosa

```
In [13]: data.describe()

Out[13]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	
count	150.000000	150.000000	150.000000	150.000000	150.000000	
mean	75.500000	5.843333	3.054000	3.758667	1.198667	
std	43.445368	0.828066	0.433594	1.764420	0.763161	
min	1.000000	4.300000	2.000000	1.000000	0.100000	
25%	38.250000	5.100000	2.800000	1.600000	0.300000	
50%	75.500000	5.800000	3.000000	4.350000	1.300000	
75%	112.750000	6.400000	3.300000	5.100000	1.800000	
max	150.000000	7.900000	4.400000	6.900000	2.500000	

```
In [15]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  --
0   Id           150 non-null    int64
1   SepalLengthCm 150 non-null    float64
2   SepalWidthCm  150 non-null    float64
3   PetalLengthCm 150 non-null    float64
4   PetalWidthCm  150 non-null    float64
5   Species       150 non-null    object
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

handling missing values

```
In [17]: data.isnull().sum()

Out[17]:
Id                0
SepalLengthCm    0
SepalWidthCm     0
PetalLengthCm    0
PetalWidthCm     0
Species          0
dtype: int64
```

handling categorical data

```
In [21]: data.drop(columns='Id', inplace=True)

In [ ]: # performing train test

In [43]: if 'Id' in data.columns:
data = df.drop(columns='Id')
le = LabelEncoder()
data['Species'] = le.fit_transform(data['Species'])

In [45]: X = data.drop(columns='Species')
y = data['Species']

In [47]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

In [49]: nn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)

Out[49]:
KNeighborsClassifier
KNeighborsClassifier(n_neighbors=1)

In [51]: y_pred = knn.predict(X_test)

In [53]: print("Accuracy:", accuracy_score(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred, target_names=le.classes_))

Accuracy: 1.0

Classification Report:
              precision    recall  f1-score   support


Iris-setosa      1.00      1.00      1.00        10
Iris-versicolor  1.00      1.00      1.00         9
Iris-virginica   1.00      1.00      1.00        11

   accuracy
   macro avg      1.00      1.00      1.00        30
   weighted avg      1.00      1.00      1.00        30
```

for different k values

```
In [55]: for k in k_values:
knn = KNeighborsClassifier(n_neighbors=k)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
acc = accuracy_score(y_test, y_pred)
accuracies.append(acc)

In [57]: # Plotting the results
plt.figure(figsize=(10, 6))
plt.plot(k_values, accuracies, marker='o')
plt.title('KNN Accuracy for Different K Values')
plt.xlabel('Number of Neighbors (K)')
plt.ylabel('Accuracy')
plt.xticks(k_values)
plt.grid(True)
plt.show()
```




The plot shows how the classification accuracy changes with different values of K. Typically, lower values of K can lead to overfitting, while higher values can result in underfitting.

Evaluate model using accuracy, confusion matrix.

```
In [66]: from sklearn.metrics import confusion_matrix, classification_report
best_k = 5
knn = KNeighborsClassifier(n_neighbors=best_k)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

In [68]: accuracy = accuracy_score(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)
report = classification_report(y_test, y_pred, target_names=le.classes_)

In [70]: plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fnc='d', cmap='Blues',
            xticklabels=le.classes_, yticklabels=le.classes_)
plt.title('Confusion Matrix (K=best_k)')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.tight_layout()
plt.show()
```



```
In [72]: accuracy, report

Out[72]: (1.0,
{'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 10, 'Iris-setosa': 1.0, 'Iris-versicolor': 1.0, 'Iris-virginica': 1.0},
{'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 9, 'Iris-setosa': 0.0, 'Iris-versicolor': 1.0, 'Iris-virginica': 0.0},
{'precision': 1.0, 'recall': 1.0, 'f1-score': 1.0, 'support': 11, 'Iris-setosa': 0.0, 'Iris-versicolor': 0.0, 'Iris-virginica': 1.0})
```

With K = 5, the model achieved: Accuracy: 100% Confusion Matrix: All predictions were correct for each class.

Classification Report: Precision, recall, and F1-score are all 1.00 across all classes.

Visualize decision boundaries.

```
In [100]: # Use only the first two features
X = data.iloc[:, :2].values
y = data['Species'].values

In [102]: # Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


In [104]: knn = KNeighborsClassifier(n_neighbors=5)
knn.fit(X_train, y_train)

# Create mesh grid for decision boundaries
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                    np.arange(y_min, y_max, 0.01))


In [106]: Z = knn.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

In [107]: # Plot
plt.figure(figsize=(10, 6))
plt.contourf(xx, yy, Z, alpha=0.4, cmap=plt.cm.coolwarm)
scatter = plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.coolwarm, edgecolor='black')
plt.xlabel('SepalLengthCm')
plt.ylabel('SepalWidthCm')
plt.title('KNN Decision Boundaries (First Two Features)')

Out[107]: Text(0.5, 1.0, 'KNN Decision Boundaries (First Two Features)')
```



```
In [108]: species_names = le.classes_
class_ids = np.unique(y)
handles = [plt.Line2D([1], [], marker='o', color='w', label=species_names[i],
                    markerfacecolor=plt.cm.coolwarm(i / 2), markersize=10) for i in class_ids]
plt.legend(handles=handles, title='Species')
plt.show()
```



In []: